

CS60038: Assignment 2

Building a simple file system with an emulated in-memory disk.

Assignment Date: 29 October 2020

Deadline: 20 November 2020

The objective of this assignment is to get hands-on experience in building a simple file system.

This assignment will involve building three parts:

- **PART A: Building an in-memory disk emulator**
- **PART B: Implementing a simple file system - SFS**
- **PART C: Implementing directory structure**

PART A: Disk Emulator

Since it is difficult to test our file system on a physical raw disk, we will need to emulate it. For that we are going to create an in-memory disk emulator using that will store blocks of data in the memory. This will provide APIs for block level access to the emulated disk. Along with that, the disk will maintain some stats about its usage.

Our disk emulator interface will have the following functions:

```
int create_disk(disk *diskptr, int nbytes);

int read_block(disk *diskptr, int blocknr, void *block_data);

int write_block(disk *diskptr, int blocknr, void *block_data);

int free_disk(disk *diskptr);
```

The disk will have fixed block size of `4KB` (4096 bytes). In addition to providing the API for reading and writing blocks, it will maintain some metadata about the disk.

```
uint32_t size; // size of the disk
uint32_t blocks; // number of usable blocks (except stat block)
```

```
uint32_t reads; // number of block reads performed
uint32_t writes; // number of block writes performed
char **block_arr; // pointer to the disk blocks
```

These stats will take up 24 bytes of storage at the beginning of the disk.

Therefore, if we consider a disk of size 409600 bytes = (4KB x 100) , the number of usable blocks will be 99 . Again, if we consider a disk of size 409624 bytes = (4KB x 100) + 24 , the number of usable blocks will be 100 .

Implement the disk emulator functions as follows:

disk * create_disk(int nbytes);

Creates a disk of size nbytes . Initializes disk struct with calculated number of blocks, and reads, writes as 0. Returns the pointer to the created disk (disk *) if successful, NULL for error.

You have to store the disk as an array of blocks. The pointer to the array of blocks, as well as other metadata will be kept in the disk structure:

```
typedef struct disk {
    uint32_t size; // size of the disk
    uint32_t blocks; // number of usable blocks (except stat block)
    uint32_t reads; // number of block reads performed
    uint32_t writes; // number of block writes performed
    char **block_arr; // array (of size blocks) of pointers to 4KB blocks
} disk;
```

int read_block(disk *diskptr, int blocknr, void *block_data)

Assumes that block_data is a 4KB buffer. Checks if blocknr is a valid block pointer - blocknr ranges from 0 to (B -1) , where B is the total number of blocks in the disk. Reads the block contents into the buffer. Returns 0 if successful, -1 otherwise. Also updates the diskptr to increment the reads count.

int write_block(disk *diskptr, int blocknr, void *block_data)

Assumes that the block_data is a 4KB buffer. Checks if blocknr is a valid block pointer. Writes the contents of the block_data buffer into the correct block. Returns 0 if success, -1 otherwise. Also, updates the diskptr to increment the writes counter.

NOTE: Updating disk_stat does not affect the writes counter.

```
int free_disk(disk *diskptr);
```

Free the memory including all blocks.

PART B: Simple File System SFS

Using the emulated disk we will be implementing a simple file system: SFS. Let us first take a look at the file system layout:

SFS will have the following components:

- Super Block
- Inodes and Inode Blocks
- Data Blocks
- Indirect Blocks
- Inode Bitmap
- Data block bitmap

Super Block

The first block of SFS will be the super block. The super block will have the following info:

```
typedef struct super_block {
    uint32_t magic_number; // File system magic number
    uint32_t blocks;       // Number of blocks in file system (except
super block)

    uint32_t inode_blocks; // Number of blocks reserved for inodes == 10%
of Blocks
    uint32_t inodes;       // Number of inodes in file system == length
of inode bit map
    uint32_t inode_bitmap_block_idx; // Block Number of the first inode
bit map block
    uint32_t inode_block_idx;       // Block Number of the first inode
block

    uint32_t data_block_bitmap_idx; // Block number of the first data
bitmap block
    uint32_t data_block_idx;       // Block number of the first data
block
    uint32_t data_blocks; // Number of blocks reserved as data blocks
} super_block;
```

Inodes and Inode Blocks

```
typedef struct inode {
    uint32_t valid; // 0 if invalid
    uint32_t size; // logical size of the file
    uint32_t direct[5]; // direct data block pointer
    uint32_t indirect; // indirect pointer
} inode;
```

As we can see the size of each inode is 32 bytes . Therefore each inode block will contain 128 inodes.

The number of blocks assigned for inodes will be 10% of the total number of available blocks in the disk (except super block).

According to the total number of inodes, the length of the inode bit map is set.

Rest of the blocks are used for data blocks and bit map for data block.

Data Blocks

Data blocks are the blocks used for storing actual file contents. We can also use these blocks as indirect blocks for storing indirect block pointers. Later we will be using these data blocks for storing directory structures also.

Indirect Blocks

These are data blocks storing pointers to other data blocks. In case direct pointers are not enough to accomodate all the block pointers, the indirect blocks will be used to store more pointers. Each pointer is an integer of 4 bytes , so one indirect block will hold 1024 such pointers. Only one indirect block is allowed per inode.

Inode Bitmap

The inode bitmap is used to determine which inodes are occupied and which are free. Accordingly while creating an inode, the first free inode index is determined from this bitmap.

Data block bitmap

Similar to inode bitmap, the data block bitmap is used to find out which blocks are free or used.

The following functions need to be implemented for SFS:

```
int format(disk *diskptr);

int mount(disk *diskptr);

int create_file();

int remove_file(int inumber);

int stat(int inumber);

int read_i(int inumber, char *data, int length, int offset);

int write_i(int inumber, char *data, int length, int offset);
```

int format(disk *diskptr)

This function takes input the pointer to `disk`. In order to write the super block first we need to calculate all the values of blocks, indexes etc.. So the organization of the file system needs to be decided first. This is done as follows:

If the disk has `N` blocks (usable blocks as specified in `diskptr->blocks`), one block is reserved for the super block. Therefore we have `M = N - 1` blocks remaining. Among these `M` blocks, 10 percent is reserved for inode blocks. Therefore the number of inode blocks is `I = 0.1 * M` (take floor). Therefore the total number of inodes will be `I * 128`, which is also the length of the inode bitmap (in bits). Therefore the number of blocks reserved for the inode bitmap will be `IB = (I * 128) / (8 * 4096)` (take ceil).

Therefore the remaining number of blocks is `R = M - I - IB`. These will be used as data blocks and data block bitmap blocks. For simplicity, consider that the length of the data block bitmap is `R bits`. Therefore the number of blocks required for data block bitmap is `DBB = R / (8 * 4096)` (take ceil). Therefore, the remaining are data blocks: `DB = R - DBB`. NOTE: we only consider the first `DB` number of bits of the data block bitmap.

The values of the super block structure will be set as follows:

```
magic_number = 12345
blocks = M
inode_blocks = I
inodes = I * 128
```

```
inode_bitmap_block_idx = 1
inode_block_idx = 1 + IB + DBB
data_block_bitmap_idx = 1 + IB
data_block_idx = 1 + IB + DBB + I
data_blocks = DB
```

The super block is then written to the disk.

The bitmaps are initialized. The inodes are also initialized with `valid = 0` . All these are written to the disk.

Returns 0 if successful, -1 otherwise.

```
int mount(disk *diskptr)
```

This method first attempts to read the super block and then verify the `magic_number` . If this is successful, a valid SFS is detected. Return 0 if successful, -1 otherwise. The except `format` , and `mount` , other functions should return error when operating on unmounted filesystem. (you need to maintain the state i.e the mounted disk pointer)

```
int create_file()
```

Creates a new inode and returns the inode pointer - that is the index of the inode (starting from 0). Initialize the size to 0 and valid to 1. Accordingly update the inode bitmap. Return -1 if error or if no inodes are left.

```
int remove_file(int inumber)
```

Removes an inode along with all corresponding data blocks and indirect pointer blocks. Just set valid to 0 in the inode, and update the bitmaps (both inode and data block bitmaps).(No need to overwrite data blocks or indirect pointer blocks).

```
int stat(int inumber)
```

Returns the stats of the input inode: logical size, and number of data blocks in use, number of direct and indirect pointers in use.

```
int read_i(int inumber, char *data, int length, int offset)
```

First checks if `inumber` , `offset` and `length` are valid or not. In case the `inumber` is invalid or `offset` is out of range return error (`-1`). If length is out of range, read only till the end of the file and return the length of the data read (in bytes). The data is read into the buffer `data` , which is assumed to be of length `length` . In case of any error return `-1` .

```
int write_i(int inumber, char *data, int length, int offset)
```

First checks if `inumber` and `offset` are valid. Then write data to a valid inode by copying `length` bytes from the buffer `data` into the data blocks of the inode starting at offset bytes. Allocate any necessary direct and indirect blocks in the process. Also update the bitmaps accordingly. Returns the number of bytes actually written. In case of any error return `-1` . If the disk is full (no free data blocks left, write as much data as possible and return the number of bytes written).

PART C: File and directory structure in SFS

Now since we have the basic functions to create, remove, read and write files, we need build a directory structure to organize files. Use the functions above to create your own directory structure. Remember every directory is also a file, and the file contains a list of its children. The list should contain:

- valid bit (unset when a file or directory is deleted and can be replace by some other entry later)
- type bit: file / directory
- the name of the file or directory (fix the allowed maximum length)
- length of the name
- inumber of the file or directory

Assign the first inode for the root directory file.

Implement the following additional functions to read and write files with file path:

```
int read_file(char *filepath, char *data, int length, int offset);
int write_file(char *filepath, char *data, int length, int offset);
int create_dir(char *dirpath);
int remove_dir(char *dirpath);
```

NOTE: In all the functions, take special care to free any allocated memory that is not required once the function returns.

For bitmap manipulations you may follow this tutorial:

<http://www.mathcs.emory.edu/~cheung/Courses/255/Syllabus/1-C-intro/bit-array.html>

Important Information

The students have to submit the source code compatible with the header files provided. A Makefile is also included for compiling the code. Implement the disk emulator code and SFS code in disc.c and sfs.c respectively.

In every source code file, students need to comment their name, roll number.

Do not make any assumption. If you find anything ambiguous in the assignment statement, do not hesitate to post it to teams.

DO NOT COPY CODE FROM OTHER GROUPS. WE WILL NOT EVALUATE A SOLUTION IF FOUND COPIED. ALL INVOLVED GROUPS WILL BE AWARDED ZERO.